

1 UNITED STATES PATENT APPLICATION

2

3 TITLE OF THE INVENTION

4 MULTI-AGENT ARTIFICIAL INTELLIGENCE SYSTEM FOR DISCOVERY, ANALYSIS,
5 GOVERNANCE, AND PARETO-PRIORITIZATION OF NOVEL, HIGH-IMPACT
6 QUESTIONS

7

8 FIELD OF THE INVENTION

9 The invention relates to machine-learning architectures and, more particularly, to distributed multi-
10 agent systems that **pro-actively surface “unknown-unknown” questions**, evaluate them on multi-
11 objective criteria, and output an **auditable, Pareto-optimal frontier** of inquiries for strategic,
12 scientific, ethical, or philosophical exploration.

13

14 BACKGROUND OF THE INVENTION

15 Progress in every discipline is bounded not by the answers we possess, but by the **questions we have**
16 **not yet imagined**. Existing AI tools excel at (i) retrieving answers to *known* queries and (ii)
17 generating questions only as a by-product of answering tasks. These tools lack:

- 18 1. a specialized engine for **detecting conceptual white-space** across heterogeneous corpora;
19 2. a **multi-agent adversarial/co-operative loop** capable of scoring each new question on
20 *novelty, impact, feasibility, ethical risk, and cross-domain leverage*;
21 3. a **transparent governance layer** that can quarantine bio-security or dual-use hazards *before*
22 public disclosure; and
23 4. an **immutable, regulator-ready audit trail** that supports reproducibility and compliance.

24 Consequently, enterprises waste resources exploring redundant or low-impact avenues, while
25 transformative research questions remain undiscovered.

26

27 **SUMMARY OF THE INVENTION**

28 The invention remedies these limitations through a **five-component architecture**:

- 29 1. **Question Discovery Agents (QDAs)** generate candidate inquiries by contrasting predictive
30 gaps in source corpora with anomaly signals from unsupervised models.
- 31 2. **Question Analysis Agents (QAAs)** compute a *score vector* $\langle \text{novelty, strategic-impact,}$
32 *feasibility, ethical-risk, cross-domain-leverage* \rangle for each candidate.
- 33 3. **Question Governance Agents (QGAs)** enforce policy, resolve scoring disputes via
34 confidence-weighted voting, and **quarantine questions whose ethical-risk exceeds a**
35 **programmable threshold**.
- 36 4. A **Question Ledger**—an append-only, cryptographically signed record—maintains full
37 provenance, including agent rationales.
- 38 5. A **Priority Engine** performs a **multi-objective optimization** over the score vectors,
39 generating a **Pareto frontier** and publishing a rank-ordered queue tailored to user-defined
40 weights.

41 Key technical advantages:

- 42 1. **True novelty detection** via adversarial generation–analysis loops plus anomaly metrics.
- 43 2. **Policy-aware quarantine workflow** that no prior art discloses.
- 44 3. **Explainability module** that attaches a chain-of-thought digest to every prioritized question.
- 45 4. **Horizontal scalability** through micro-service deployment of agent instances.

46

BRIEF DESCRIPTION OF THE DRAWING

FIG. 1 (see page 2 of the specification) shows the high-level data-flow among Data Ingestion (101), QDAs (102), QAAs (104), QGAs (105), Question Ledger & Priority Engine (106–108), and the User/Down-stream AI interface (109). Solid arrows denote primary data flow; dashed arrows show governance feedback.

DETAILED DESCRIPTION OF THE INVENTION

Priority Engine Algorithm

Upon receiving adjudicated score vectors S_i , the engine solves:

$$\max_{Q_i \in Ledger} f(S_i, W)$$

Subject to:

1. f implements a Pareto frontier search such that no selected question is dominated on all objectives.
2. W is a user-supplied weight vector enabling dynamic re-ranking without re-analysis.

Quarantine Workflow

If:

$$ethical_risk(Q_i) > \theta$$

(where θ is a policy-defined threshold), the Question Governance Agent (QGA) moves Q_i to a **quarantine sub-ledger**.

Release is subject to:

1. Applicable human oversight policies
2. A revised risk assessment procedure

74 **EXAMPLE USE CASES**

- 75 • **Synthetic-biology red-team:** The system auto-quarantines a question whose lab-protocol
76 implications exceed a BSL-3 risk score; human bio-safety officers review and approve partial
77 disclosure.
- 78 • **Corporate strategy:** Pareto frontier highlights three cross-domain R&D opportunities that
79 traditional pipeline tools missed, saving 18 months of exploratory spend.
- 80 • **National-security foresight:** High-risk cyber-warfare questions are flagged and routed to
81 cleared analysts under sealed audit keys.